Cached-candidate refinement diagnostic

64 score-selected modes from an incomplete, partial-cache source run. Snapshot dated 2026-10-26.

The refined score winner remains wrong by 853.4 km. This diagnostic does not establish a localization fix or any release gate.

Top result by refined score The two bars track the refined score winner from its source seed to its refined result; this seed was not the source run global winner. The lowest-error alternative is tabulated separately and is not the estimate.

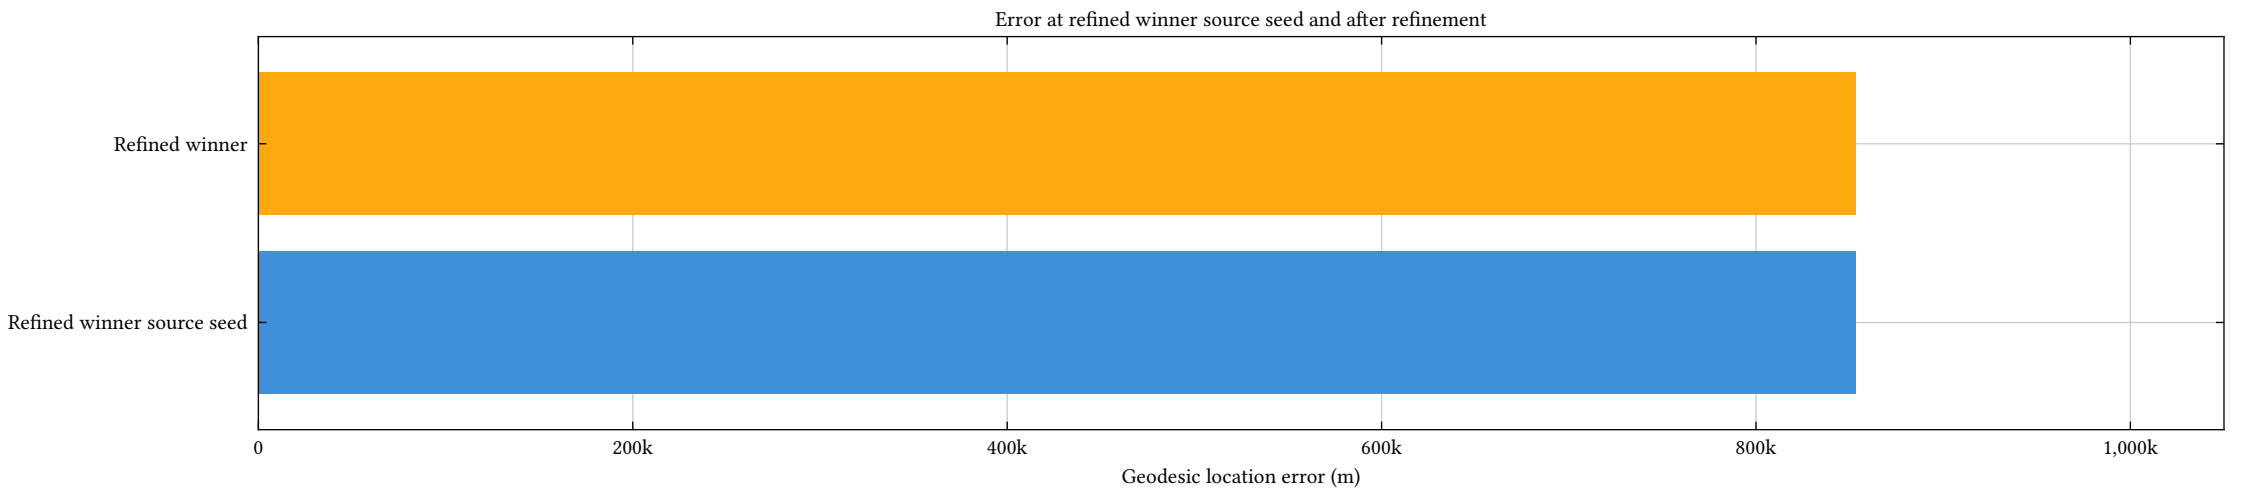

Mode role	Source rank	Refined rank	Source score	Refined score	Source to refined error (m)	True assignments
Score winner	3	1	15.238	10.324	853406.1 ⟶ 853439.8	0 / 4
Lowest-error alternative. Not estimate.	5	11	16.356	14.86	296.3 ⟶ 266.4	4 / 4

Score-ranked alternatives in the selected subset The horizontal axis is source objective rank among 47,182 exact-scored hypotheses. The 64 displayed modes were selected by search score from that incomplete run, not a complete search ranking. Paired markers show each mode before and after refinement.

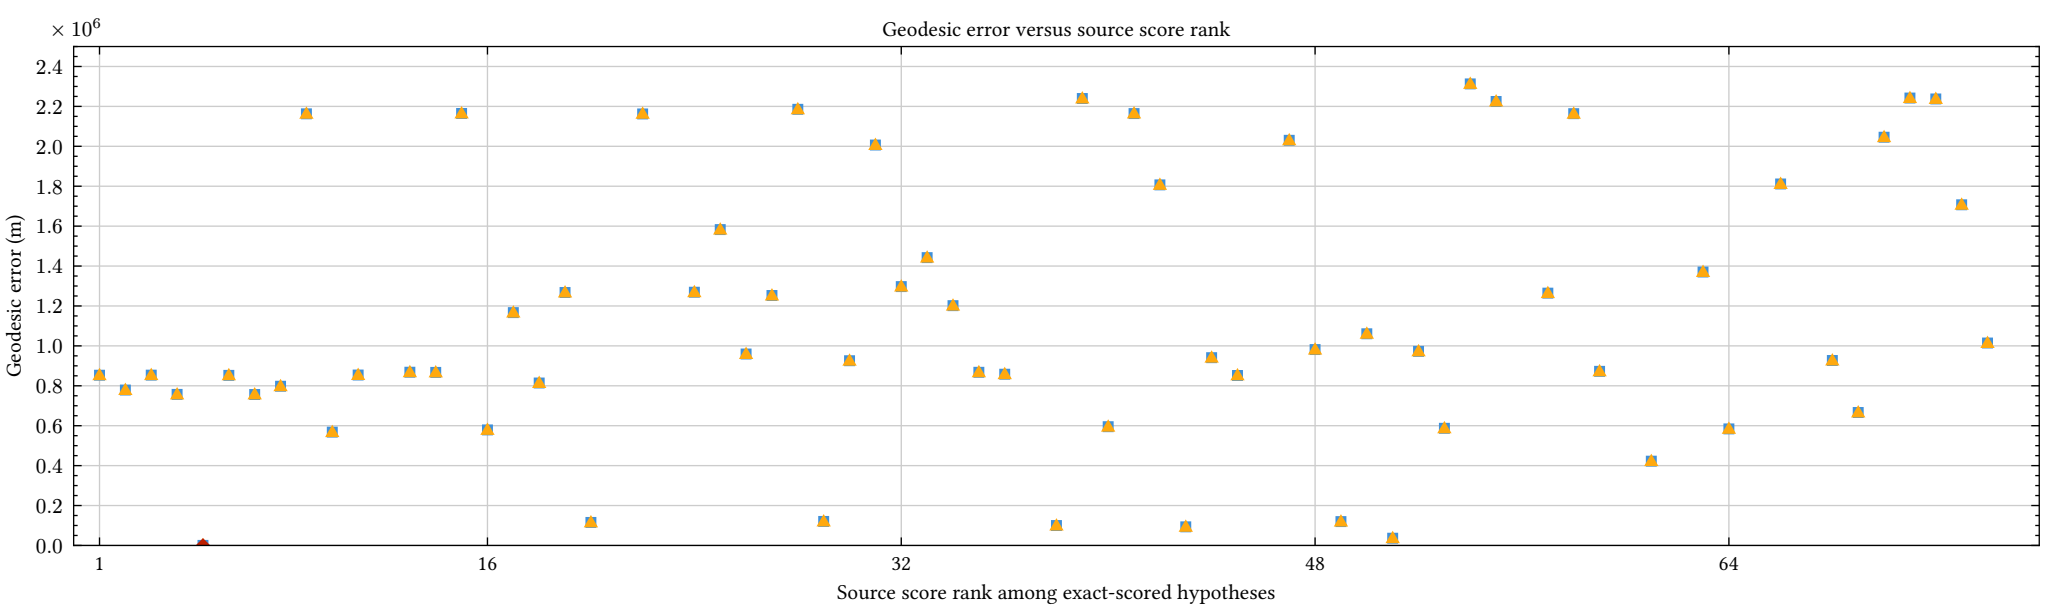

Default series order and marker key. First square is source, second triangle is refined, third diamond is the lowest-error alternative, which is not the score-selected estimate.

Refinement lowered objective score in 64 / 64 modes; location error improved in 37 and worsened in 27. Median geodesic error changed from 950060.5 m to 950085.8 m. The source run had 47182 exact-scored hypotheses, 82 unresolved hypotheses, and 96.8 percent of its candidate population in an unknown state at interruption. Evaluator artifact SHA-256: 88290e7ac580c0cd7bd32fb7fd813e0ff7e55ac3a7aefbcca5ee295f0a0a79cd. Input manifest SHA-256: 1cc41596e787c990950b4d1f00043c73b4815ffaf9e2371c762bf356d3bba063. Evaluator source commit: 624e0469eb2131e301e77d07a7f727a9b80cd063. Truth was used only for independent evaluation. Authorship J.C. Vaught.